



JOB → STAGES → TASKS



.collect()



Job #1



⋮

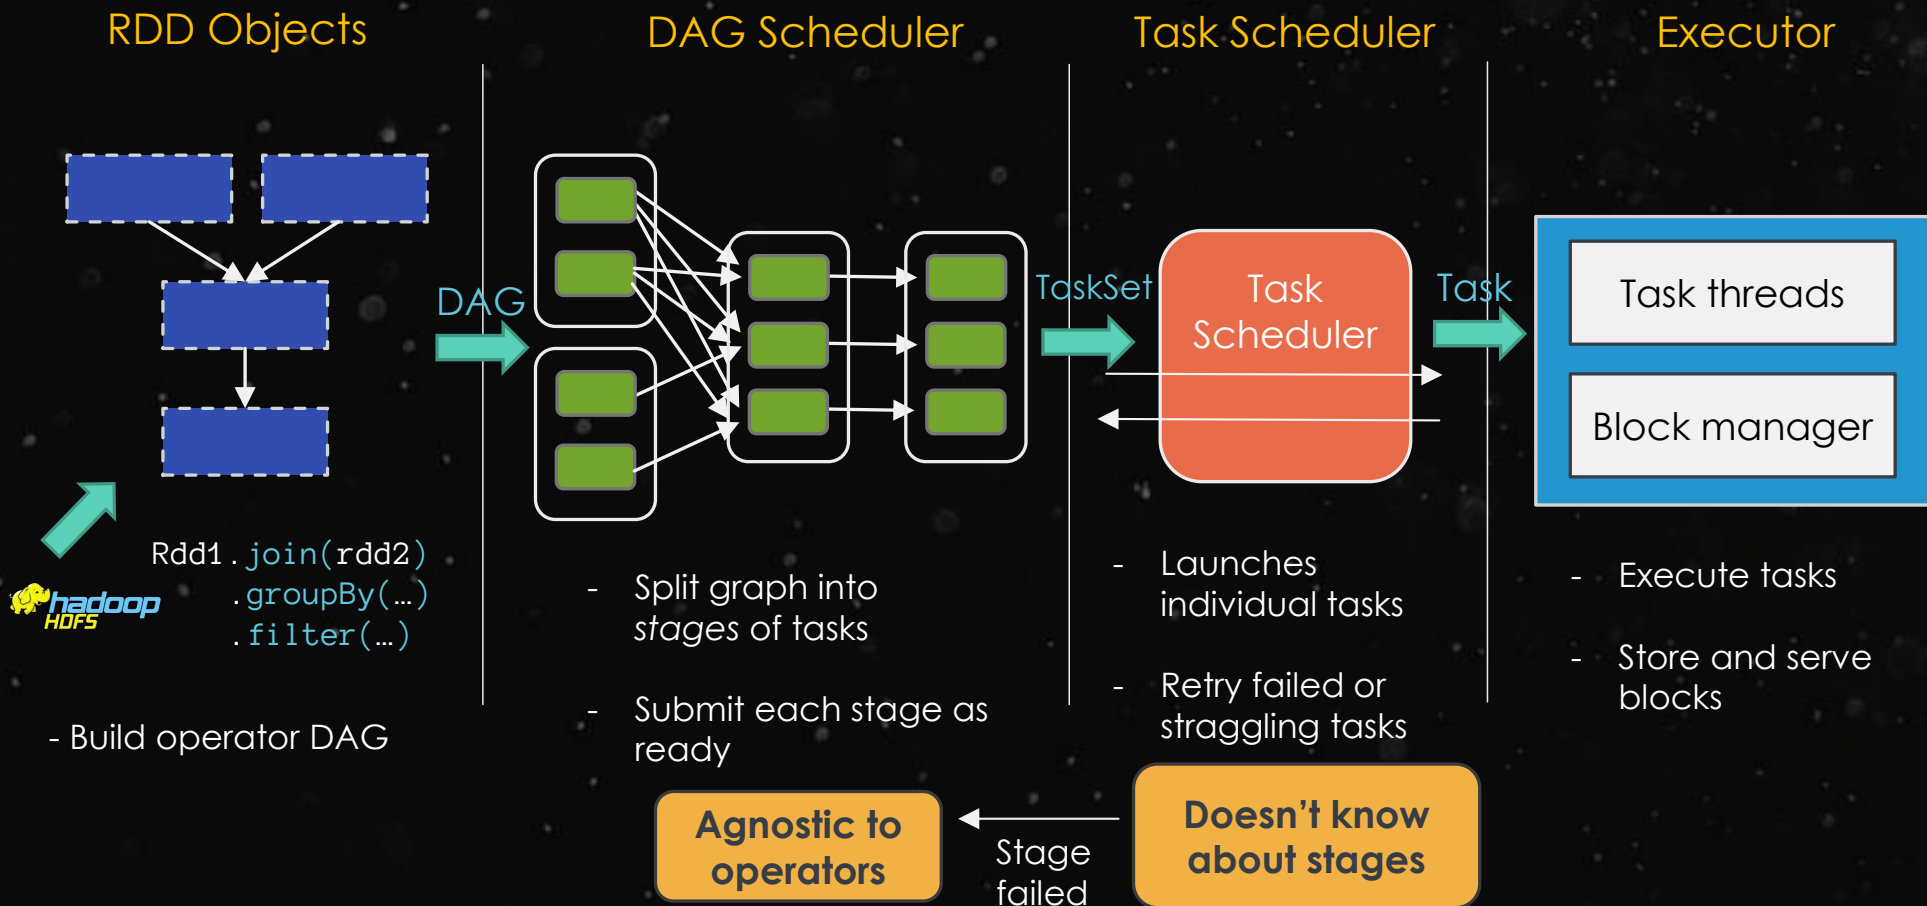


Task #1
Task #2
Task #3

⋮



SCHEDULING PROCESS



LINEAGE

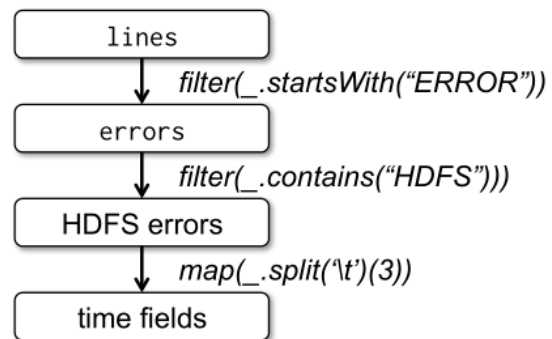


Figure 1: Lineage graph for the third query in our example. Boxes represent RDDs and arrows represent transformations.

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
errors.persist()
```



LINEAGE

“One of the challenges in providing RDDs as an abstraction is choosing a representation for them that can track lineage across a wide range of transformations.”

“The most interesting question in designing this interface is how to represent dependencies between RDDs.”


“We found it both sufficient and useful to classify dependencies into two types:

- narrow dependencies, where each partition of the parent RDD is used by at most one partition of the child RDD
- wide dependencies, where multiple child partitions may depend on it.”




LINEAGE DEPENDENCIES

Narrow Dependencies:



Wide Dependencies:




Requires
shuffle


Examples of narrow and wide dependencies.


Each box is an RDD, with partitions shown as shaded rectangles.

STAGES

 = RDD

 = cached partition

 = lost partition



LINEAGE

Dependencies: **Narrow** vs **Wide**

"This distinction is useful for two reasons:

1) Narrow dependencies allow for pipelined execution on one cluster node, which can compute all the parent partitions. For example, one can apply a map followed by a filter on an element-by-element basis.

In contrast, wide dependencies require data from all parent partitions to be available and to be shuffled across the nodes using a MapReduce-like operation.

2) Recovery after a node failure is more efficient with a narrow dependency, as only the lost parent partitions need to be recomputed, and they can be recomputed in parallel on different nodes. In contrast, in a lineage graph with wide dependencies, a single failed node might cause the loss of some partition from all the ancestors of an RDD, requiring a complete re-execution."



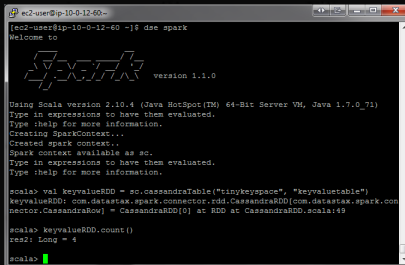
To display the lineage of an RDD, Spark provides a `toDebugString` method:

```
scala> input.toDebugString
```

```
res85: String =  
(2) data.text MappedRDD[292] at textFile at <console>:13  
| data.text HadoopRDD[291] at textFile at <console>:13
```

```
scala> counts.toDebugString
```

```
res84: String =  
(2) ShuffledRDD[296] at reduceByKey at <console>:17  
+- (2) MappedRDD[295] at map at <console>:17  
| FilteredRDD[294] at filter at <console>:15  
| MappedRDD[293] at map at <console>:15  
| data.text MappedRDD[292] at textFile at <console>:13  
| data.text HadoopRDD[291] at textFile at <console>:13
```





How do you know if a shuffle will be called on a Transformation?

- repartition, join, cogroup, and any of the *By or *ByKey transformations can result in shuffles
- If you declare a numPartitions parameter, it'll probably shuffle
- If a transformation constructs a shuffledRDD, it'll probably shuffle
- combineByKey calls a shuffle (so do other transformations like groupByKey, which actually end up calling combineByKey)

Note that repartition just calls coalesce w/ True:

```
RDD.scala    def repartition(numPartitions: Int)(implicit  
              ord: Ordering[T] = null): RDD[T] = {  
                coalesce(numPartitions, shuffle = true)  
              }
```



How do you know if a shuffle will be called on a Transformation?

Transformations that use “numPartitions” like `distinct` will probably shuffle:

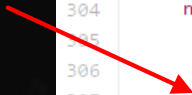
```
def distinct(numPartitions: Int)(implicit ord: Ordering[T] =  
  null): RDD[T] =  
  map(x => (x, null)).reduceByKey((x, y) => x,  
    numPartitions).map(_._1)
```

PERSERVES PARTITIONING


- An extra parameter you can pass a k/v transformation to let Spark know that you will not be messing with the keys at all
- All operations that shuffle data over network will benefit from partitioning
- Operations that benefit from partitioning:
cogroup, groupWith, join, leftOuterJoin, rightOuterJoin, groupByKey, reduceByKey, combineByKey, lookup, . . .

<https://github.com/apache/spark/blob/master/core/src/main/scala/org/apache/spark/rdd/RDD.scala#L302>

```
299  /**
300   * Return a new RDD containing only the elements that satisfy a predicate.
301   */
302  def filter(f: T => Boolean): RDD[T] = {
303    val cleanF = sc.clean(f)
304    new MapPartitionsRDD[T, T](
305      this,
306      (context, pid, iter) => iter.filter(cleanF),
307      preservesPartitioning = true)
308  }
```



[Support](#)[Developers](#)[Contact Us](#)[Downloads](#)


Ask Bigger Questions

[COMMUNITY](#)[DOCUMENTATION](#)[DOWNLOADS](#)[TRAINING](#)[BLOGS](#)

Hadoop & Big Data

Our Customers

FAQs

Blog

Accumulo (1)

Avro (16)

Bigtop (6)

Books (11)

Careers (14)

CDH (150)

Cloud (20)

Cloudera Labs (4)

Cloudera Life (6)

Cloudera Manager (73)

Community (208)

How-to: Tune Your Apache Spark Jobs (Part 1)

by Sandy Ryza | March 09, 2015 | no comments

Learn techniques for tuning your Apache Spark jobs for optimal efficiency.

(Editor's note: Sandy presents on "[Estimating Financial Risk with Spark](#)" at Spark Summit East on March 18.)

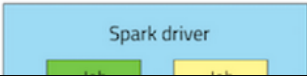
When you write Apache Spark code and page through the public APIs, you come across words like *transformation*, *action*, and *RDD*. Understanding Spark at this level is vital for writing Spark programs. Similarly, when things start to fail, or when you venture into the web UI to try to understand why your application is taking so long, you're confronted with a new vocabulary of words like *job*, *stage*, and *task*. Understanding Spark at this level is vital for writing *good* Spark programs, and of course by *good*, I mean *fast*. To write a Spark program that will execute efficiently, it is very, very helpful to understand Spark's underlying execution model.

In this post, you'll learn the basics of how Spark programs are actually executed on a cluster. Then, you'll get some practical recommendations about what Spark's execution model means for writing efficient programs.

How Spark Executes Your Program

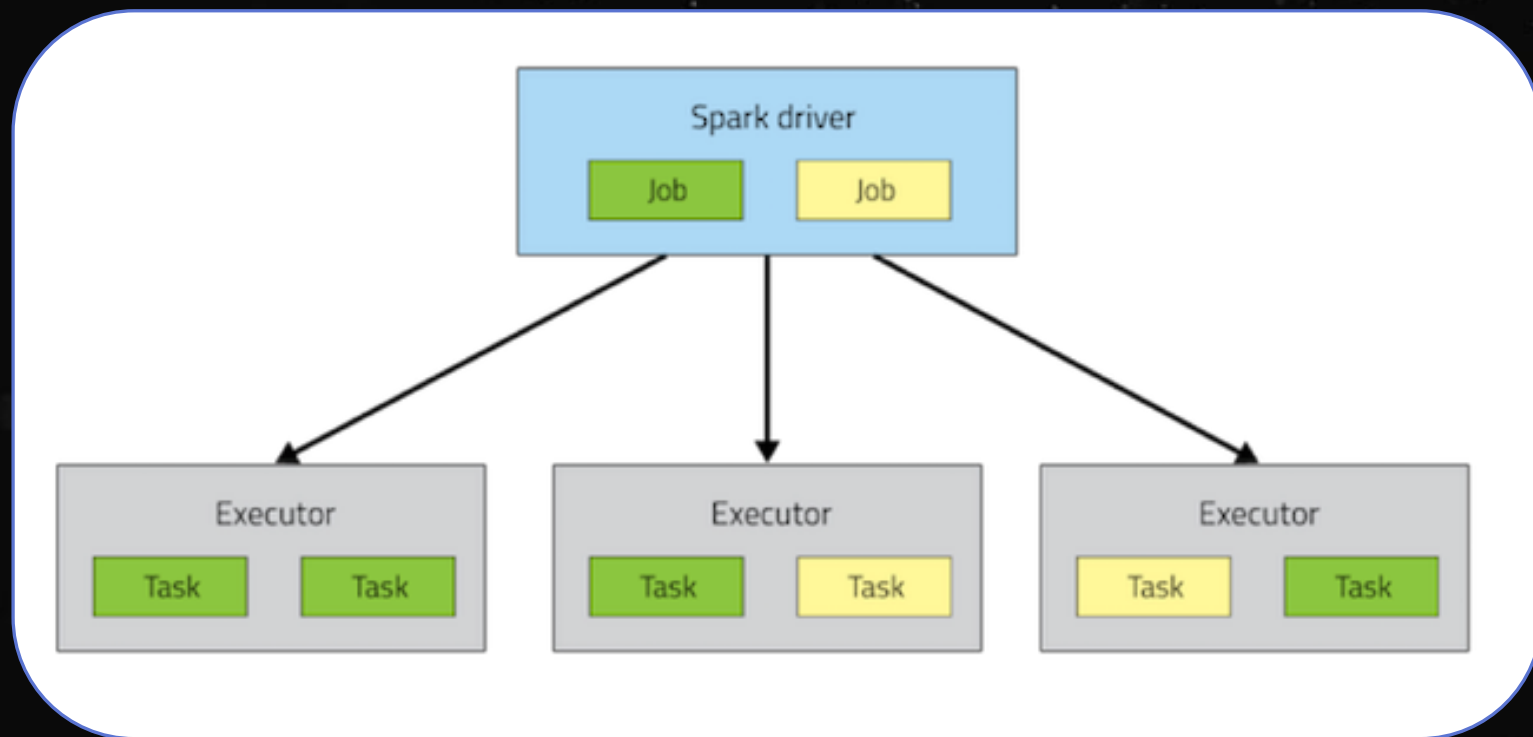
A Spark application consists of a single *driver* process and a set of *executor* processes scattered across nodes on the cluster.

The driver is the process that is in charge of the high-level control flow of work that needs to be done. The executor processes are responsible for executing this work, in the form of *tasks*, as well as for storing any data that the user chooses to cache. Both the driver and the executors typically stick around for the entire time the application is running, although [dynamic resource allocation](#) changes that for the latter. A single executor has a number of slots for running tasks, and will run many concurrently throughout its lifetime. Deploying these processes on the cluster is up to the cluster manager in use (YARN, Mesos, or Spark Standalone), but the driver and executor themselves exist in every Spark application.



[Link](#)

Source: Cloudera



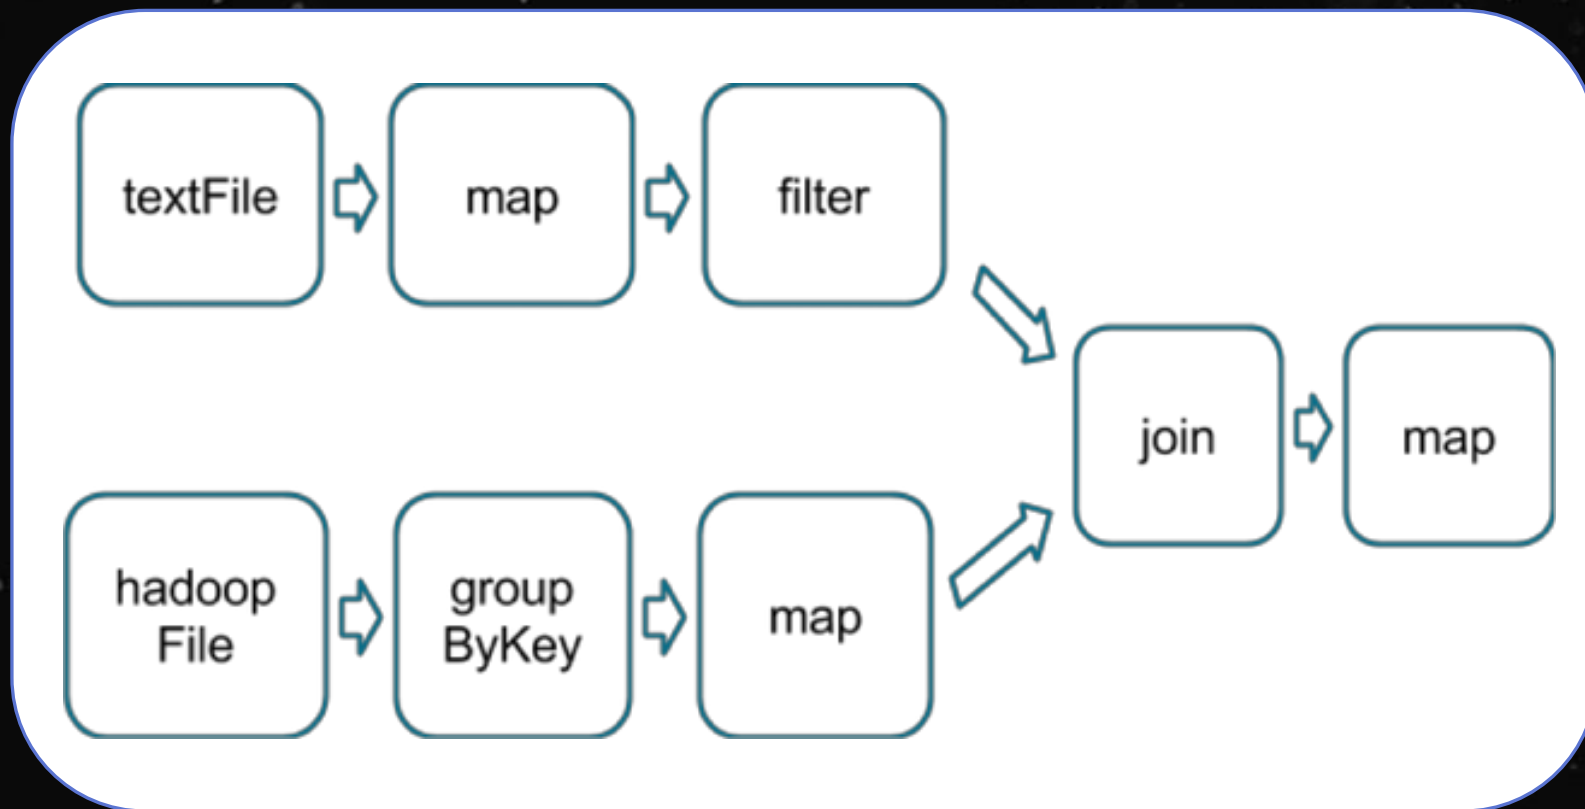
Source: Cloudera

How many Stages will this code require?

```
sc.textFile("someFile.txt").  
  map(mapFunc).  
  flatMap(flatMapFunc).  
  filter(filterFunc).  
  count()
```

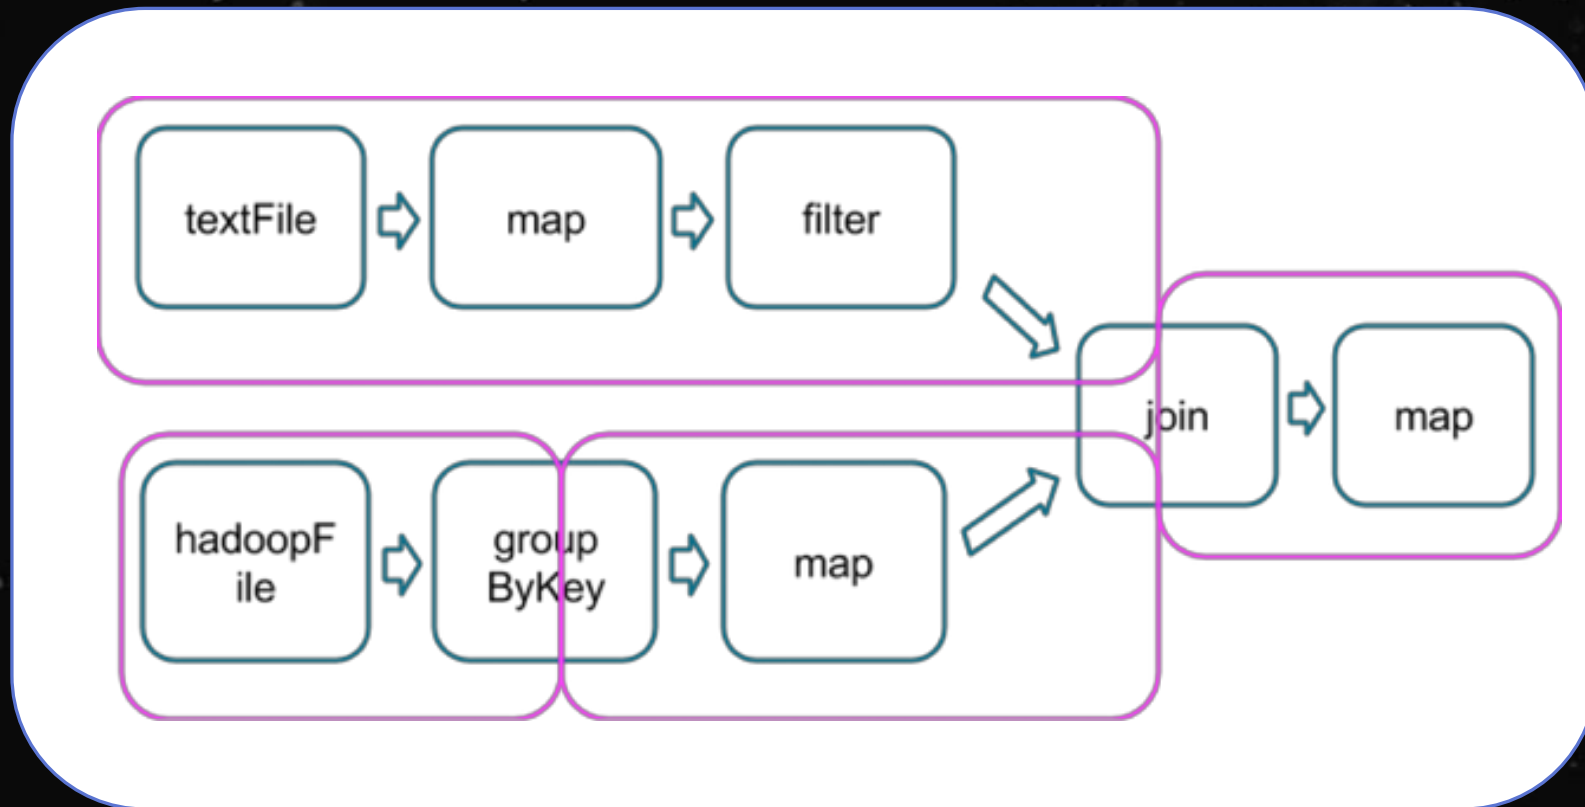
Source: Cloudera

How many Stages will this DAG require?



Source: Cloudera

How many Stages will this DAG require?



Source: Cloudera